

# Unikernels: Library Operating Systems for the Cloud

Anil Madhavapeddy, Richard Mortier, Charalampos Rotsos, David Scott, Balraj Singh, Thomas Gazagnaire, Steven Smith, Steven Hand, Jon Crowcroft

Jay Sh  
COMP 517  
2020.10.12



# Unikernels: The Rise of the Virtual Library Operating System

Anil Madhavapeddy, David J. Scott

Jay Sh  
COMP 517  
2020.10.12

A decorative light blue triangle is located in the bottom right corner of the slide.

# Problem & Motivation

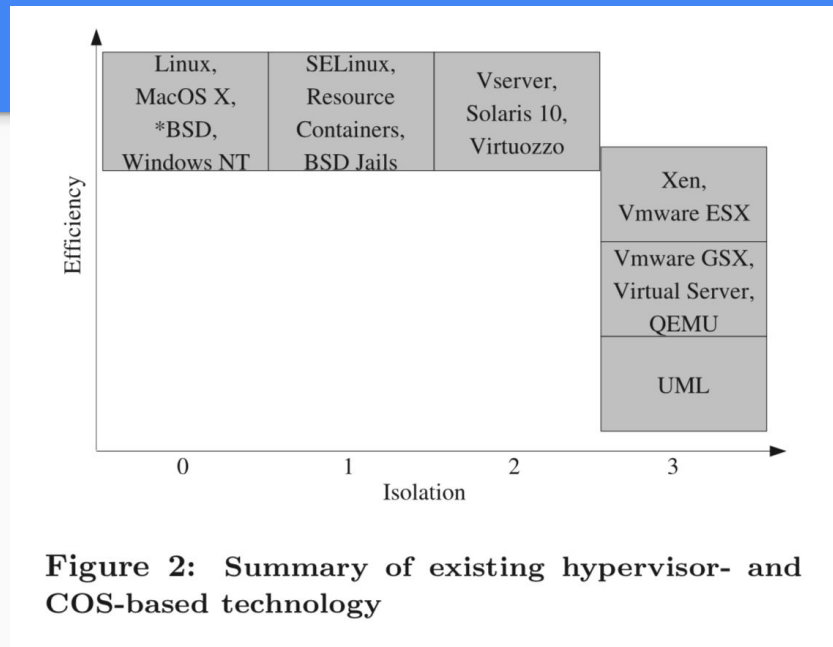
- VMs are used extensively in the cloud to achieve economy of scale
  - Xen, KVM
- VMs are often single purpose
  - Commodity OSs are not
- The hosting software stack has too many layers
  - Physical protocols, POSIX, processes and threads, language runtimes, application code
- VM images often need to be transferred for deployment
  - Commodity OSs are not optimized for size or boot time

# Problem & Motivation

- libOS is good for applications that need performance but has two problems
  - Resource isolation is hard when running multiple applications
  - Device drivers need to be rewritten and catch up with commodity hardware
- “The fast-moving world of commodity PC hardware meant that, no matter how many graduate students were tasked to write drivers, any research libOS prototype was doomed to become obsolete in a few short years.”

# Container-Based OSs?

- Relaxed isolation
- Cannot optimize system components to suit application needs
  - Networking stack



Container-based Operating System Virtualization: A Scalable, High-performance Alternative to Hypervisors

# Unikernel

- Virtual machine running single language libOS and application
- OS components are implemented as libraries and applications link against them
- **Single process, single address space**
- Designed for the cloud

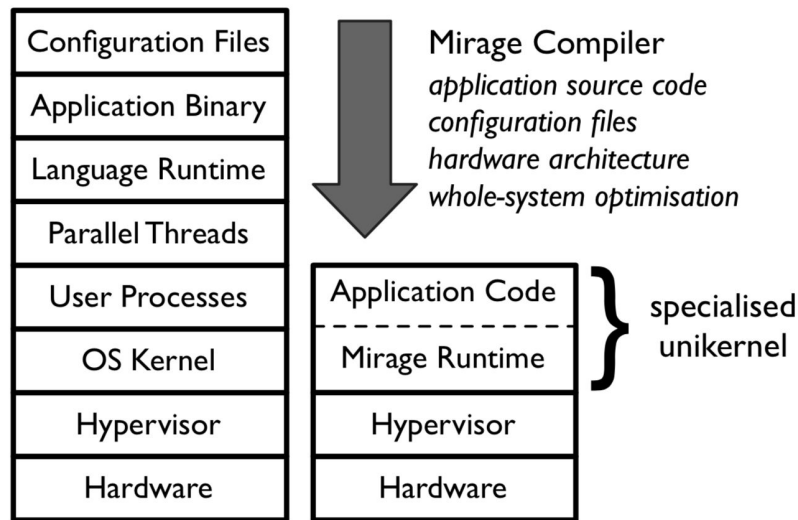


Figure 1: Contrasting software layers in existing VM appliances vs. unikernel's standalone kernel compilation approach.

# MirageOS

- Functional prototype that runs on Xen
- libOS and application written in OCaml
- Uses PVBoot library to boot unikernels
  - One vCPU
  - One 64-bit address space
- Modified OCaml runtime
  - Occupies text and data region
  - GC uses two heaps
- Uses Lwt library for threading

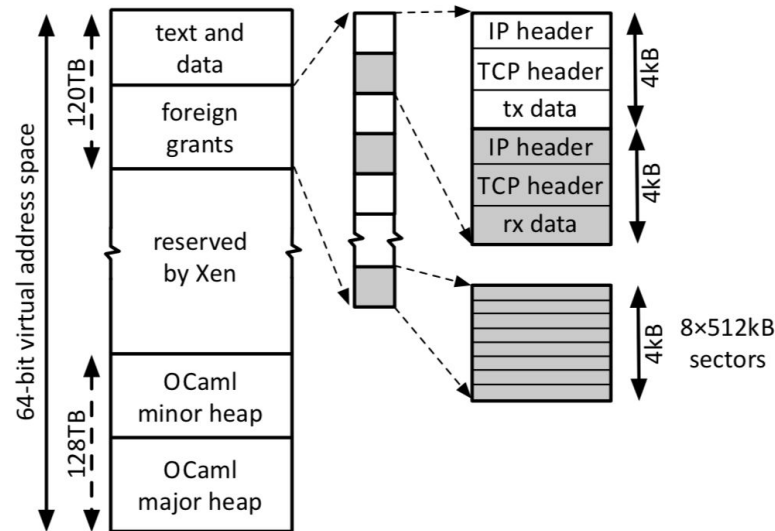


Figure 2: Specialised virtual memory layout of a 64-bit Mirage unikernel running on Xen.

# Device I/O

- All device drivers are written in OCaml
- Xen devices are manipulated by writing to and reading from a shared memory page of slots
  - OCaml is used to safely access these pages
- Zero-copy I/O
  - Two VMs share a grant table that maps a page to an offset in the table that can be used to reference and access the page from another VM
  - OCaml is used to track and free contents of the grant table
    - High order functions that free resources when the wrapped function terminates or errors out



# Protocol I/O


- Uses vchan for inter VM communication
  - Pages of shared buffer that other VMs can access
- Applications link to network protocol libraries that send and receive packets by reading and writing I/O pages as described before
- Same for block devices, also using I/O pages
  - Applications link to filesystem and caching libraries that can be customized

# Security

- Thread model
  - Hypervisor is trusted
  - Applications trust external entities via protocols such as SSL
- Stack canaries
  - Special values placed after a buffer and monitored for modification
- Guard pages
  - Unallocated pages after allocated pages that cause segfaults when accessed
- Sealing
  - When booting, mark pages as either writable or executable but not both
  - Asks the hypervisor to disallow page table modification
- Compile time address space randomization
  - When compiling, randomly shuffle the positions of heap, stack, ...

# Security

- Static type checking and garbage collected heap prevent memory errors
- Only system components needed are linked = smaller attack surface



(b) Key components of the Mirage unikernel.

# Evaluation • Boot Time

- Unikernels boot fast

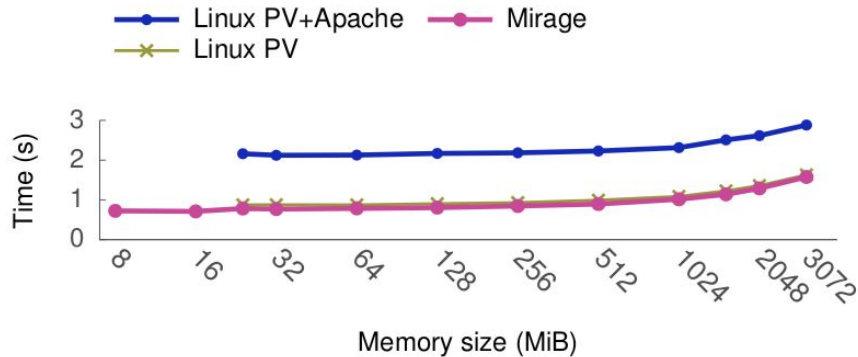


Figure 5: Domain boot time comparison.





Figure 6: Boot time using an asynchronous Xen toolstack.

# Evaluation • Threading


- Threads are faster to create due to garbage collected heap management



(a) Creation times.

# Evaluation • Threading

- Thread timers are more accurate due to no kernel/userspace boundary



(b) Jitter for  $10^6$  parallel threads sleeping and waking after a fixed period.

# Evaluation • Networking

- 10% higher latency due to type safety, higher receive throughput due to no userspace copying, lower transmit throughput due to higher CPU usage

<i>Configuration</i>	<i>Throughput [ std. dev. ] (Mbps)</i>			
	1 flow		10 flows	
Linux to Linux	1590	[ 9.1 ]	1534	[ 74.2 ]
Linux to Mirage	1742	[ 18.2 ]	1710	[ 15.1 ]
Mirage to Linux	975	[ 7.0 ]	952	[ 16.8 ]

Figure 8: Comparative TCP throughput performance with all hardware offload disabled.

# Evaluation • Storage

- Comparable random block read throughput




Figure 9: Random block read throughput, +/- 1 std. dev.



# Evaluation • Web Application

- Mirage scales better




Figure 12: Simple dynamic web appliance performance.





Figure 13: Static page serving performance, comparing Mirage and Apache2 running on Linux.

# Evaluation • Size

- Fewer LoC, small binary size



(a) Key cloud components vs. the Mirage unikernel codebase.

Appliance	Binary size (MB)	
	Standard build	Dead code elimination
DNS	0.449	0.184
Web Server	0.673	0.172
OpenFlow switch	0.393	0.164
OpenFlow controller	0.392	0.168

Table 2: Sizes of Mirage unikernels, before and after dead-code elimination. Configuration and data are compiled directly into the unikernel.

# Takeaway

- Single purpose virtual machines
- libOS and application written in one language \*
- One process, one address space
- Light weight, secure, fast

**Table 2. Other unikernel implementations.**

<b>Unikernel</b>	<b>Language</b>	<b>Targets</b>
Mirage <sup>13</sup>	OCaml	Xen, kFreeBSD, POSIX, WWW/js
Drawbridge <sup>17</sup>	C	Windows "picoprocess"
HalVM <sup>8</sup>	Haskell	Xen
ErlangOnXen	Erlang	Xen
OSv <sup>2</sup>	C/Java	Xen, KVM
GUK	Java	Xen
NetBSD "rump" <sup>9</sup>	C	Xen, Linux kernel, POSIX
ClickOS <sup>14</sup>	C++	Xen

# Discussion

- What are the tradeoffs between containers and unikernels? How would you choose between containers and unikernels?
- Are you satisfied with or convinced by the security that unikernels claim to provide?
- What are other potential use cases for unikernels?
- Do you agree with writing system libraries in high level language? Why is using type safe languages important for unikernels?
- MirageOS does not provide POSIX compatibility. Is it a good idea?